Aviral Srivastava

• Github: kebab-mai-haddi • LinkedIn: /in/sriavi • Website: aviralsrivastava.com
• Email: hi@avisri.com • ▓▓▓▓▓▓▓ • Blog

EDUCATION	

EDUCATION

Boston University, Boston, MA.   Sep 2019 - Dec 2020 (Expected)
- MS, Computer Science. CGPA : 3.77/4
  - **Subjects**: Designing Data Warehouse, Database Security, Advanced Database Management Systems, Graduate Distributed Systems (MIT), ML, AI, Adv. Health Info and Big Data, Advanced Analysis of Algorithms, Information Systems Analysis and Design.

Vellore Institute of Technology, Chennai, India.   Jul 2014 - May 2018
- B.Tech, Computer Science and Engineering. CGPA : 3.56/4

SKILLS

- **Languages**: Python, C++, Go, Java, PHP, Rust, SQL, MySQL, PostgreSQL
- **Technologies**: Spark, Dask, Kubernetes, Docker, Vagrant, Argo, Airflow, Kafka, Pulsar, Celery, Flink, Snowflake, Cassandra, RocksDB, PrestoDB, KyotoDB, Redis, Django.
- **Platforms**: AWS, GCP, Azure

WORK EXPERIENCE

Twitter, Software Engineer   Apr 2021 - Present
- As part of the Graph Storage team, I am working on real-time storage infrastructure.
- Led the team for building a managed graph database offering for analytics purposes,graph machine learning, and applications. Reduced onboarding time from 1 week to minutes, for teams at Twitter.

Redhat, Software Engineering (Data) Intern, Internal Data-Hub team, Boston   May 2020 - Present
- Re-partitioned the schema of SOS reports (**10 Trillion+ rows**, **100+ TeraBytes** data) to ensure high data availability and low latency. **Reduced** the count of S3 objects by more than **90%**: from 15M+ objects to < 1M objects.
- Scaling the visualization dashboards for internal analytical teams to prevent outages and provide smooth user journey for our data team members.
- Standardized data workflows for the internal teams using **Argo, Spark, JupyterHub, Kubernetes**

Boston University, Research Associate (Software Engineering), Boston   Sep 2019 - May 2020
- Got funded by MIT-ICORPS for commercializing the research for improving software quality.

Atlan, Software Engineer(Data)   May 2018 - Jul 2019
- Atlan is a data democratization startup that helps data teams collaborate frictionlessly on data projects.
- Ensured the development and on-premise deployments of dashboards like United Nations SDGs monitoring, Ujjwala Dashboard using **Python3, Kubernetes, Docker and AWS S3**.
- Integrated serverless JupyterHub: a multi-user hub that spawns, manages, and proxies multiple instances of the single-user Jupyter notebook server. Stack: **papermill, Python3, Docker, Helm, Kubernetes**. Used by companies like **Unilever, Mindshare, McKinsey & Co.**
- Built and deployed Data Repository: a version control system for tabular data, including features like 'git diff', rollback, at a scale of over **billion rows in one cycle**. Stack: **Python3, S3, RocksDB, KyotoDB, Hive metastore and PrestoDB**. This scaled the production releases by **300%**: 1 project release to 3 projects release per quarter.
- Built a micro-service with Event Sourcing mechanism for payments in Collect using **Kafka**.

Centify, Co-founder, CTO   Mar 2018 - Jul 2018
- Audience Engagement tools which enable digital media platforms to integrate social media reactions with video or textual content.
- Built the complete system for processing social media content and applying emotional as well as sentimental analysis. Got featured on IBM's developer blog

Wingify, SDE Intern   Dec 2017 - Mar 2018
- Wingify,with clients like **ebay, Disney, AMD, etc.**, has a product VWO which is an experience optimization platform.
- Reduced the query processing time for the product by **60%** by building a micro-service that parses textual inputs leveraging the Database.
- Stack: **PHP, MySQL, CICD, Docker, Vagrant, Blackbox Testing, Unit testing, JMeter**

OPEN SOURCE

- Profiling RocksDB-Apache Flink
- ZCash Service Status Dashboard
- Other projects